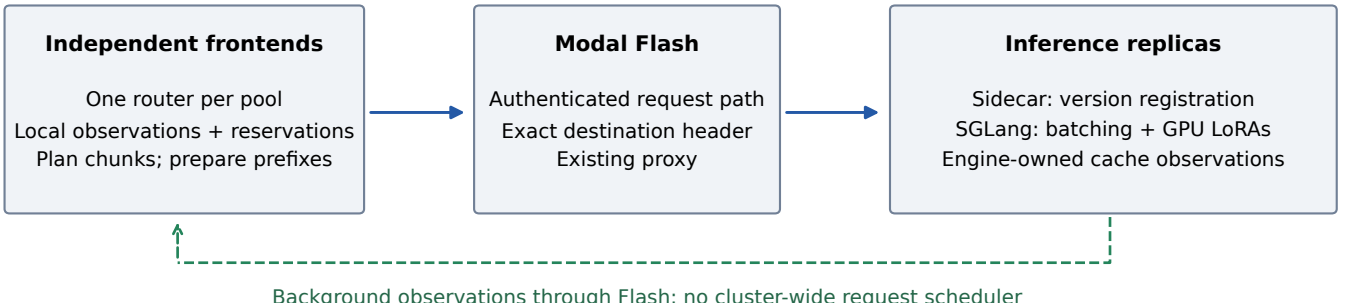
Independent frontends
One router per pool
Local observations + reservations
Plan chunks; prepare prefixes
Modal Flash
Authenticated request path
Exact destination header
Existing proxy
Inference replicas
Sidecar: version registration
SGLang: batching + GPU LoRAs
Engine-owned cache observations
Background observations through Flash; no cluster-wide request scheduler